


scanplan — three front-ends, one core

The tiers differ almost only in HOW METRIC SCALE IS RECOVERED. Everything downstream sees one intermediate representation.



scanplan run <capture> **the tier is detected from the input**
LiDAR tier uses NO neural model: plane fitting, rasterisation, morphology, connected components, least squares.
Deterministic and offline — it cannot fail at a demo because a checkpoint did not cache.